


# Multi layer perceptrons

Dr Suresh Sundaram


# Single layer perceptron



$$y = g \left( b + \sum_i x_i w_i \right)$$


A mathematical equation for the output of a single layer perceptron. The output  $y$  is shown in pink. The equation includes a bias term  $b$  and a summation over  $i$  of the product of the  $i$ 'th input  $x_i$  and the  $i$ 'th weight  $w_i$ . A red arrow points to the summation term with the label "nonlinearity". A blue arrow points to the bias term with the label "bias". A blue arrow points to the weight term with the label "i'th weight". A green arrow points to the input term with the label "i'th input".

# Limitations of single layer perceptron




## AND gate


$x_1$	$x_2$	$y$
0	0	0
0	1	0
1	0	0
1	1	1




AND		
I <sub>1</sub>	I <sub>2</sub>	out
0	0	0
0	1	0
1	0	0
1	1	1




OR		
I <sub>1</sub>	I <sub>2</sub>	out
0	0	0
0	1	1
1	0	1
1	1	1



XOR		
I <sub>1</sub>	I <sub>2</sub>	out
0	0	0
0	1	1
1	0	1
1	1	0




# Need of multi layer perceptron : XOR Gate





A	B	O	$H_1$	$H_2$
0	0	0	1	1
0	1	1	0	1
1	0	1	0	1
1	1	0	0	0

# Need of multi layer perceptron : XOR Gate



$H_1$	$H_2$	$H$
1	1	0
0	1	1
0	1	1
0	0	0



- One of the most popular methods for training such multilayer networks is based on gradient descent in error — the back propagation algorithm (or generalized delta rule), a natural extension of the perceptron algorithm.





Figure 6.3: Whereas a two-layer network classifier can only implement a linear decision boundary, given an adequate number of hidden units, three-, four- and higher-layer networks can implement arbitrary decision boundaries. The decision regions need not be convex, nor simply connected.

# Multi layer perceptron

- Simple three-layer neural network.
- This one consists of an input layer (having  $d$  input units), a hidden layer with ( $n_H$  hidden units) and an output layer (a single unit), interconnected by modifiable weights, represented by links between layers.



# Multi layer perceptron

Net Output at input to  $j^{\text{th}}$  hidden node

$$\text{net}_j = \sum_{i=1}^d x_i w_{ji} + w_{j0} = \sum_{i=0}^d x_i w_{ji} \equiv \mathbf{w}_j^t \mathbf{x},$$

Non linearity imposed by activation function  $f(\dots)$

$$y_j = f(\text{net}_j).$$

$$f(\text{net}) = \text{Sgn}(\text{net}) \equiv \begin{cases} 1 & \text{if } \text{net} \geq 0 \\ -1 & \text{if } \text{net} < 0, \end{cases}$$

This  $f()$  is sometimes called the *transfer function* or merely “nonlinearity” of a unit,

# Multi layer perceptron

Each output unit similarly computes its net activation based on the hidden unit signals as

$$net_k = \sum_{j=1}^{n_H} y_j w_{kj} + w_{k0} = \sum_{j=0}^{n_H} y_j w_{kj} = \mathbf{w}_k^t \mathbf{y},$$

where the subscript  $k$  indexes units in the output layer

$n_H$  denotes the number of hidden units

# Multi layer perceptron

Each output unit then computes the nonlinear function of its *net*, emitting

$$z_k = f(\text{net}_k).$$

---

$$g_k(\mathbf{x}) \equiv z_k = f \left( \sum_{j=1}^{n_H} w_{kj} f \left( \sum_{i=1}^d w_{ji} x_i + w_{j0} \right) + w_{k0} \right).$$


# Multi layer perceptron

$$J(\mathbf{w}) \equiv 1/2 \sum_{k=1}^c (t_k - z_k)^2 = 1/2(\mathbf{t} - \mathbf{z})^2,$$

where  $\mathbf{t}$  and  $\mathbf{z}$  are the target and the network output vectors of length  $c$ ;  $\mathbf{w}$  represents all the weights in the network.

The backpropagation learning rule is based on gradient descent. The weights are initialized with random values, and are changed in a direction that will reduce the error:

$$\Delta \mathbf{w} = -\eta \frac{\partial J}{\partial \mathbf{w}},$$




Linear Activation Function

# Activation Functions


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$




## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$




## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$






id	color
1	red
2	blue
3	green
4	blue



id	color_red	color_blue	color_green
1	1	0	0
2	0	1	0
3	0	0	1
4	0	1	0



# Multi layer perceptron

## MEAN SQUARE ERROR LOSS


$$J(\mathbf{w}) \equiv 1/2 \sum_{k=1}^c (t_k - z_k)^2 = 1/2(\mathbf{t} - \mathbf{z})^2,$$



## CROSS ENTROPY LOSS


$$L_{\text{CE}} = - \sum_{i=1}^n t_i \log(p_i), \text{ for n classes,}$$


where  $t_i$  is the truth label and  $p_i$  is the Softmax probability for the  $i^{th}$  class.


# Error / cost function











# Multi layer perceptron

$$\Delta w_{mn} = -\eta \frac{\partial J}{\partial w_{mn}},$$


where  $\eta$  is the *learning rate*, and merely indicates the relative size of the change in weights.


# Multi layer perceptron


This iterative algorithm requires taking a weight vector  $\mathbf{w}$  at iteration  $m$  and updating it as:

$$\mathbf{w}(m + 1) = \mathbf{w}(m) + \Delta\mathbf{w}(m),$$


# Backpropagation







# Gradient descent : Back propagation



$$\frac{\partial J}{\partial w_{kj}} = \frac{\partial J}{\partial \text{net}_k} \frac{\partial \text{net}_k}{\partial w_{kj}} = \delta_k \frac{\partial \text{net}_k}{\partial w_{kj}}$$

where the *sensitivity* of unit  $k$  is defined to be

$$\delta_k \equiv -\frac{\partial J}{\partial \text{net}_k},$$

and describes how the overall error changes with the unit's activation.

$$\delta_k \equiv -\frac{\partial J}{\partial \text{net}_k} = -\frac{\partial J}{\partial z_k} \frac{\partial z_k}{\partial \text{net}_k} = (t_k - z_k) f'(\text{net}_k).$$

# Gradient descent : Back propagation


$$\frac{\partial \text{net}_k}{\partial w_{kj}} = y_j.$$

Taken together, these results give the weight update (learning rule) for the hidden-to-output weights:

$$\Delta w_{kj} = \eta \delta_k y_j = \eta(t_k - z_k) f'(net_k) y_j.$$

# Gradient descent : Back propagation

The learning rule for the input-to-hidden units is more subtle,



$$\frac{\partial J}{\partial w_{ji}} = \frac{\partial J}{\partial y_j} \frac{\partial y_j}{\partial net_j} \frac{\partial net_j}{\partial w_{ji}}.$$

$$\begin{aligned}
 \frac{\partial J}{\partial y_j} &= \frac{\partial}{\partial y_j} \left[ \frac{1}{2} \sum_{k=1}^c (t_k - z_k)^2 \right] \\
 &= - \sum_{k=1}^c (t_k - z_k) \frac{\partial z_k}{\partial y_j} \\
 &= - \sum_{k=1}^c (t_k - z_k) \frac{\partial z_k}{\partial net_k} \frac{\partial net_k}{\partial y_j} \\
 &= - \sum_{k=1}^c (t_k - z_k) f'(net_k) w_{jk}.
 \end{aligned}$$

# Gradient descent : Back propagation

$$\delta_j \equiv f'(net_j) \sum_{k=1}^c w_{kj} \delta_k.$$

the sensitivity

at a hidden unit is simply the sum of the individual sensitivities at the output units weighted by the hidden-to-output weights  $w_{jk}$ , all multiplied by  $f'(net_j)$ . Thus the learning rule for the input-to-hidden weights is:

$$\Delta w_{ji} = \eta x_i \delta_j = \eta x_i f'(net_j) \sum_{k=1}^c w_{kj} \delta_k.$$

# Gradient descent : Back propagation




Figure 6.5: The sensitivity at a hidden unit is proportional to the weighted sum of the sensitivities at the output units:  $\delta_j = f'(net_j) \sum_{k=1}^c w_{kj} \delta_k$ . The output unit sensitivities are thus propagated “back” to the hidden units.

# Training methodologies

The three most useful training protocols are: stochastic, batch and on-line.

In stochastic training (or pattern training), patterns are chosen randomly from the stochastic training set, and the network weights are updated for each pattern presentation.

In batch training, all patterns are presented to the network before learning batch (weight update) takes place.

In on-line training, each pattern is presented once and on-line protocol only once; there is no use of memory for storing the patterns.

# MOMENTUM BASED GRADIENT DESCENT

$$\Delta w_{ij} = (\eta * \frac{\partial E}{\partial w_{ij}})$$


weight increment      learning rate      weight gradient

$$\Delta w_{ij} = (\eta * \frac{\partial E}{\partial w_{ij}}) + (\gamma * \Delta w_{ij}^{t-1})$$


momentum factor      weight increment,  
previous iteration

# MOMENTUM BASED GRADIENT DESCENT



SGD without momentum




SGD with momentum




# REGULARIZATION



## Training Vs. Test Set Error





# REGULARIZATION TECHNIQUES

$L_2$  regularization

$L_1$  regularization


Data Augmentation

Early stopping


Adding noise to input

Adding noise to target


# Data Augmentation




## Early stopping




# Adding noise to inputs




# Drop-out for regularization



(a) Standard Neural Net



(b) After applying dropout.




## Drop-out for regularization

We initialize all the parameters (weights) of the network and start training

For the first training instance (or mini-batch), we apply dropout resulting in the thinned network

We compute the loss and back propagate

Which parameters will we update? Only those which are active



For the second training instance (or mini-batch), we again apply dropout resulting in a different thinned network

We again compute the loss and backpropagate to the active weights

If the weight was active for both the training instances then it would have received two updates by now

If the weight was active for only one of the training instances then it would have received only one update by now