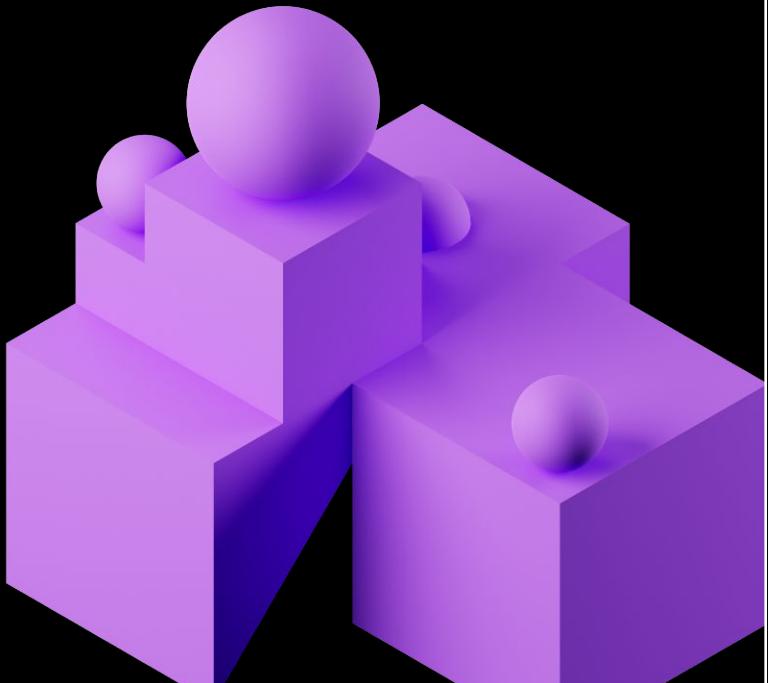




Deep Learning with Databricks

Databricks
2023

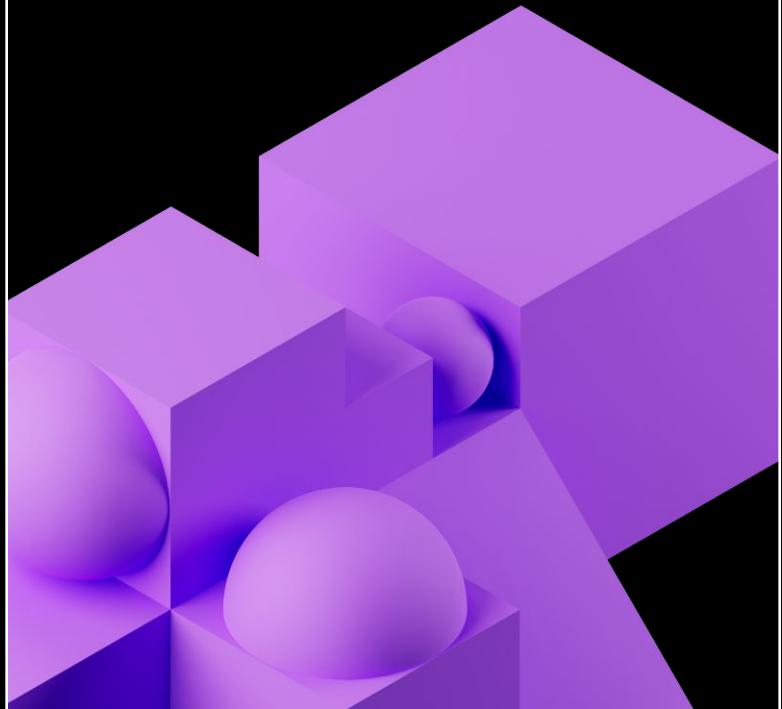


Agenda

Deep Learning with Databricks		
	Demo	Lab
Introduction		
Linear Regression with Keras	✓	✓
Callbacks	✓	✓
MLflow	✓	✓
Hyperopt	✓	✓
Horovod + Petastorm	✓	✓
Convolutional Neural Networks	✓	
Transfer Learning	✓	✓

LET'S GET STARTED

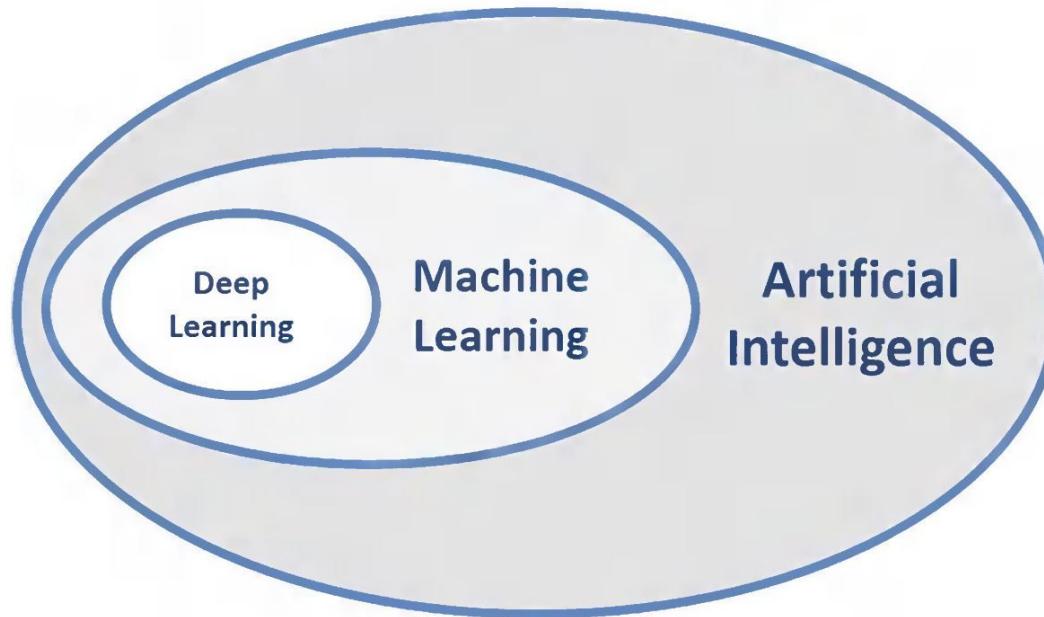
Deep Learning Overview



Why Deep Learning?

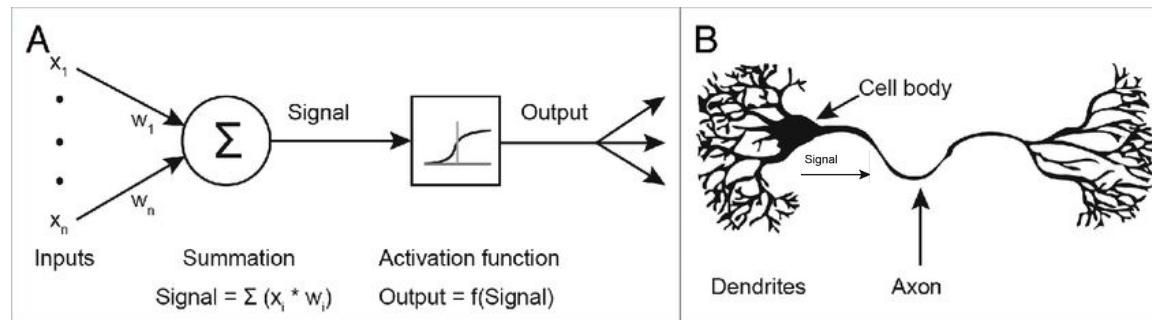
- Performs well on complex datasets like images, sequences, and natural language
- Performs better as data amount increases
- Theoretically can learn any relationship (universal approximation theorem)

Where Does DL Fit In?



What Is Deep Learning?

*Features learned from data using basic building blocks
inspired by the human brain*



A) An *artificial* neuron

B) A *biological* neuron

Keras

- High-level Python API to build neural networks
- Official high-level API of TensorFlow
- Has over 250,000 users
- Released by François Chollet in 2015
- **Sequential API** (common method) and **Functional API** (to define complex topologies)

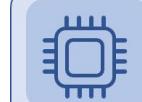
```
●●● Sequential API Example  
1 import tensorflow as tf  
2 from tensorflow.keras import datasets, layers, models  
3  
4 # Load dataset  
5 mnist = datasets.mnist  
6 (x_train, y_train), (x_test, y_test) = mnist.load_data()  
7 x_train, x_test = x_train / 255.0, x_test / 255.0  
8  
9 # Construct a neural network model  
10 model = models.Sequential()  
11 model.add(layers.Flatten(input_shape=(28, 28)))  
12 model.add(layers.Dense(512, activation=tf.nn.relu))  
13 model.add(layers.Dropout(0.2))  
14 model.add(layers.Dense(10, activation=tf.nn.softmax))  
15 model.compile(optimizer='adam',  
16                 loss='sparse_categorical_crossentropy',  
17                 metrics=['accuracy'])  
18  
19 # Train and evaluate the model  
20 model.fit(x_train, y_train, epochs=5)  
21 model.evaluate(x_test, y_test)
```

[Source](#)



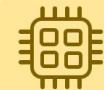
Hardware Considerations

- GPUs are often the tool of choice for DL training due to their superior parallel execution abilities over CPUs.
- CPUs are easier to program, faster for smaller datasets, and allow more flexibility of use.
- *We'll be using CPUs in this course because they are cheaper for learning.*



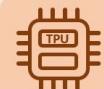
CPU

- Small models
- Small datasets
- Useful for design space exploration



GPU

- Medium-to-large models, datasets
- Image, video processing
- Application on CUDA or OpenCL



TPU

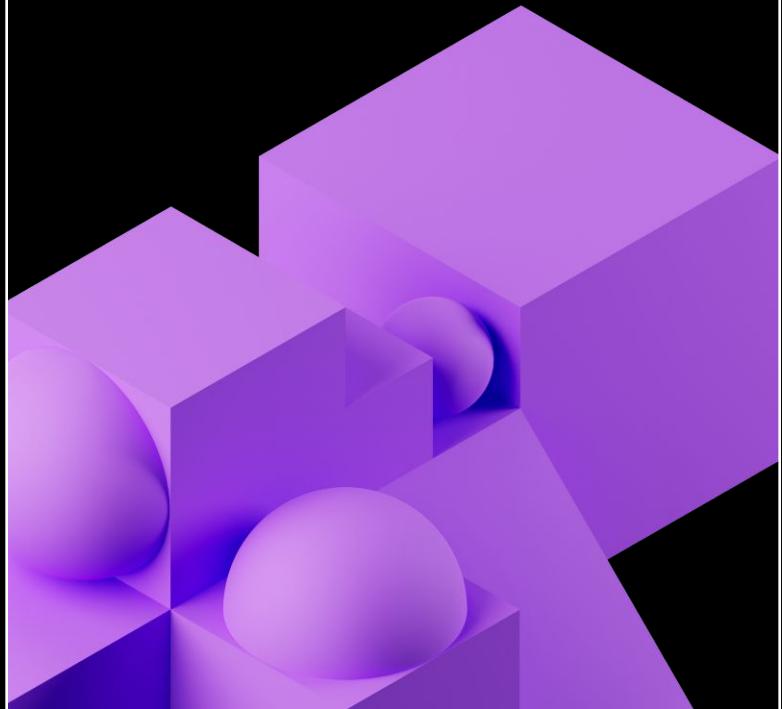
- Matrix computations
- Dense vector processing
- No custom TensorFlow operations



FPGA

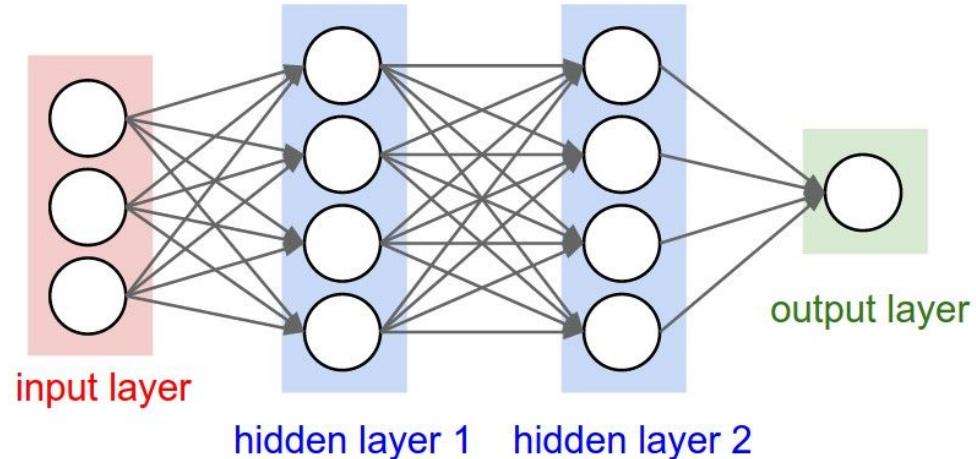
- Large datasets, models
- Compute intensive applications
- High performance, high perf./cost ratio

Neural Network Fundamentals



Layers

- Input layer
- Zero or more hidden layers
- Output layer



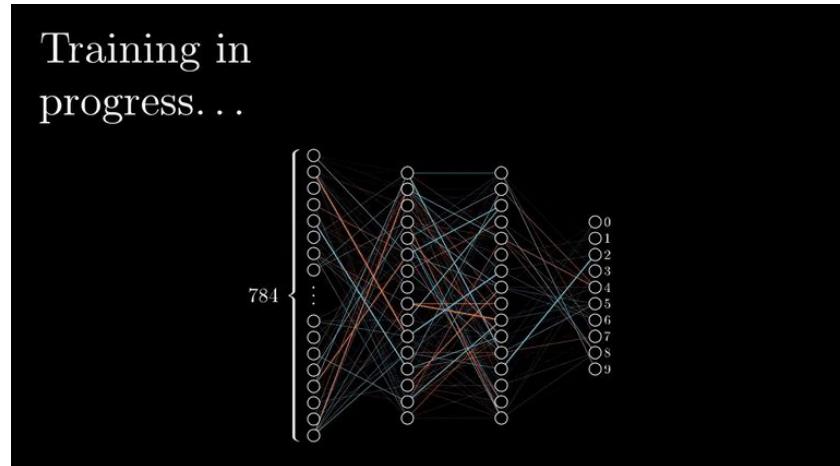
[Image source](#)



Backpropagation

While training

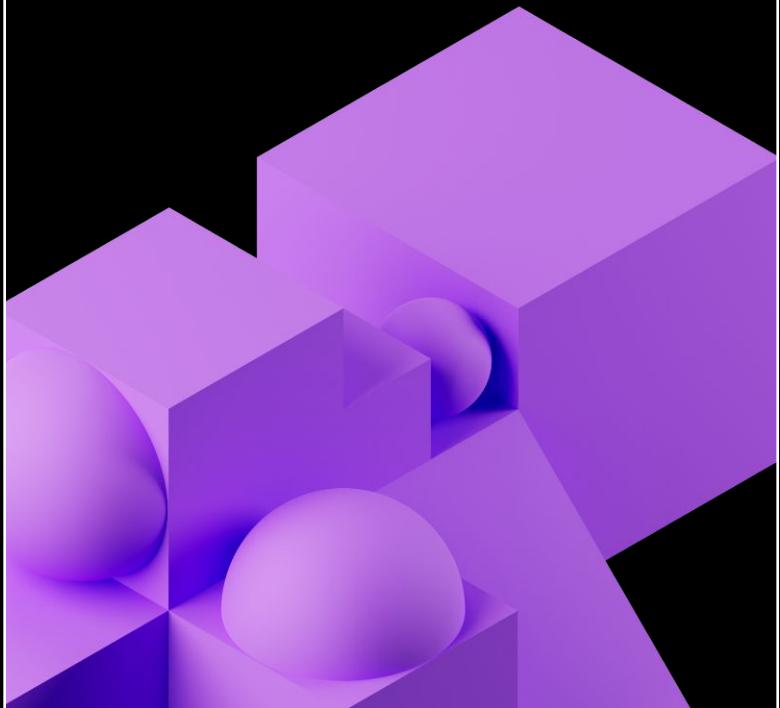
- Take an example input
- Feed it forward through the network
- Calculate the error in the prediction
- Use the error to correct the previous layers



[Image source](#)



Linear Regression with Keras



Activation Functions

Common activation functions in neural networks

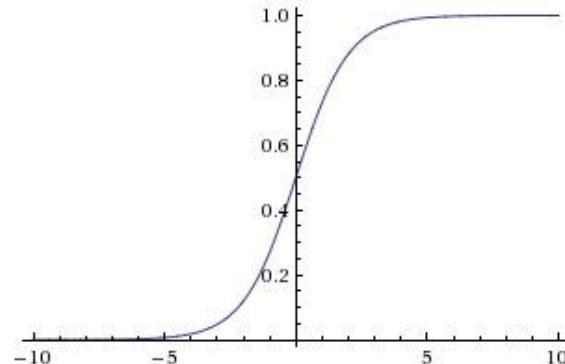
Provide non-linearity in our neural networks to learn more complex relationships

- Sigmoid
- Tangent
- ReLU (**R**ectified **L**inear **U**nit)
- Leaky ReLU
- Softmax
- ELU (**E**xponential **L**inear **U**nit)

Sigmoid

Activation function

- Saturates and kills gradients
- Not zero-centered along the y-axis
- Last layer activation function for binary classification models



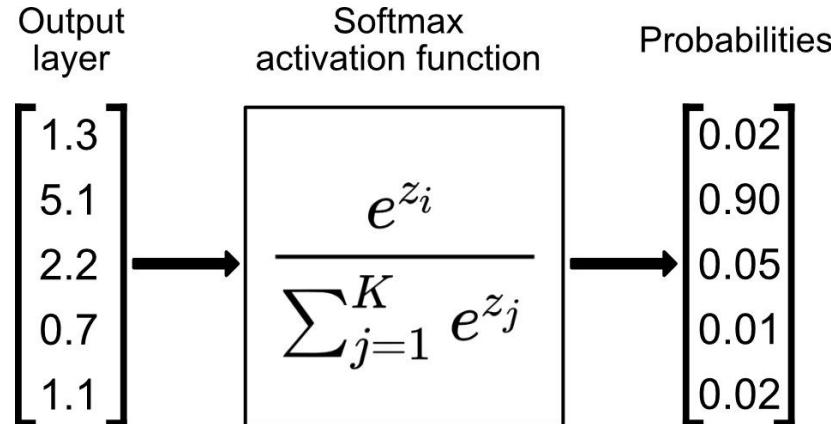
[Image source](#)



Softmax

Activation function

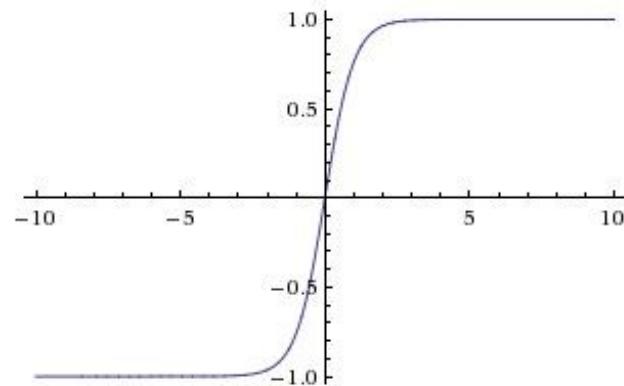
- Convert logits to probabilities
- Last layer activation function for multi-class classification



Hyperbolic Tangent (Tanh)

Activation function

- Zero centered!
- BUT, like the sigmoid, its activations saturate



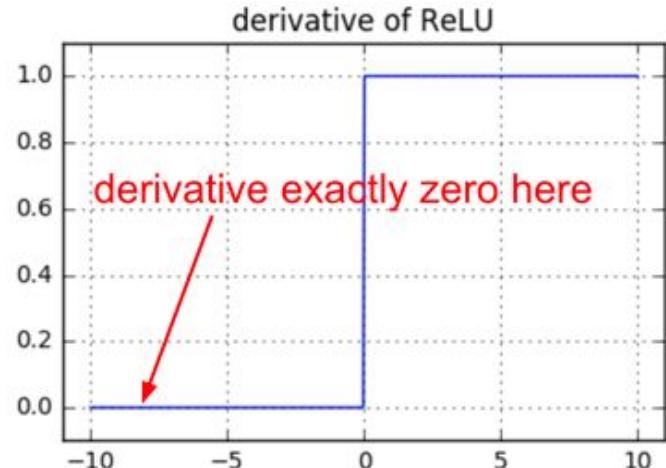
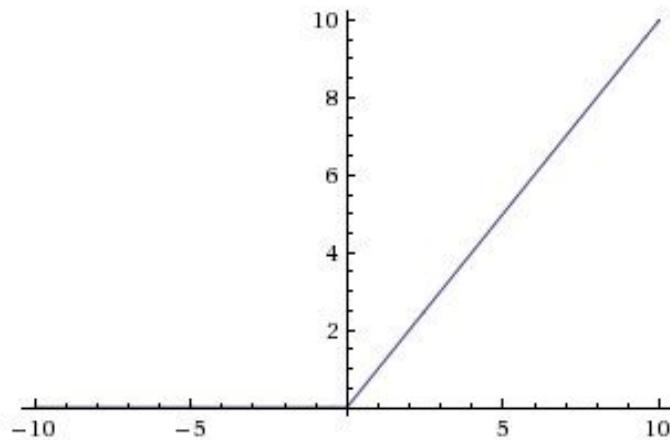
[Image source](#)



ReLU (Rectified Linear Unit)

Activation function

- BUT, gradients can still go to zero



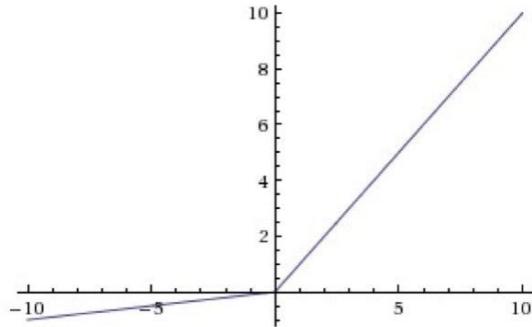
[Image source](#)



Leaky ReLU

Activation function

- For $x < 0$: $f(x) = \alpha * x$
- For $x \geq 0$: $f(x) = x$



These functions are not differentiable at 0, so we set the derivative to 0 or average of left and right derivative.

[Image source](#)

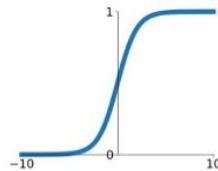


Visual Comparison

Activation Functions

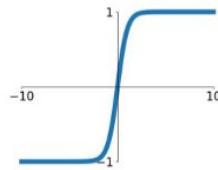
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



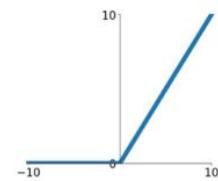
tanh

$$\tanh(x)$$



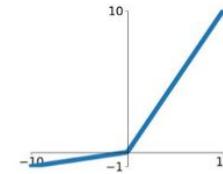
ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

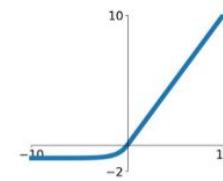


Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

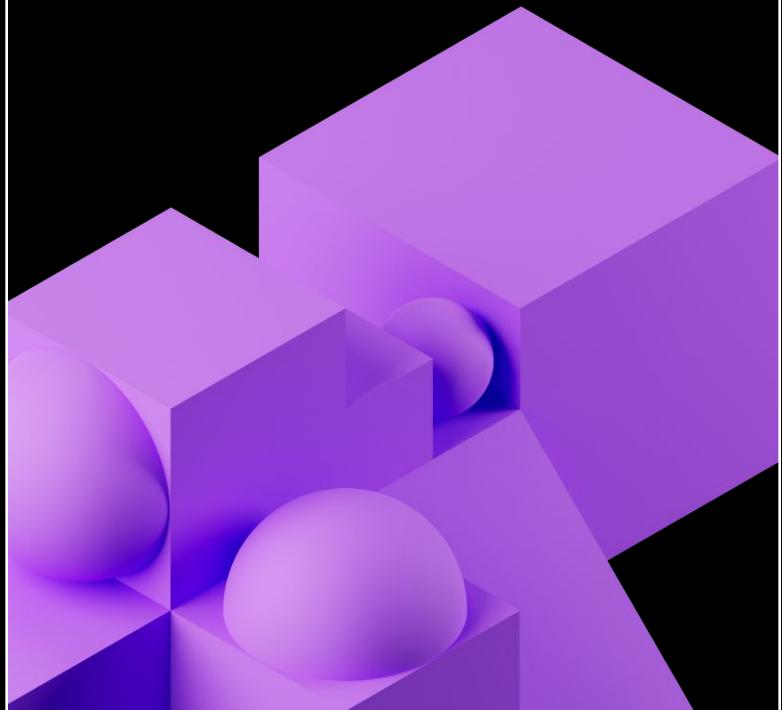
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



[Image source](#)



Optimizers



Stochastic Gradient Descent (SGD)

- Choosing a proper learning rate can be difficult

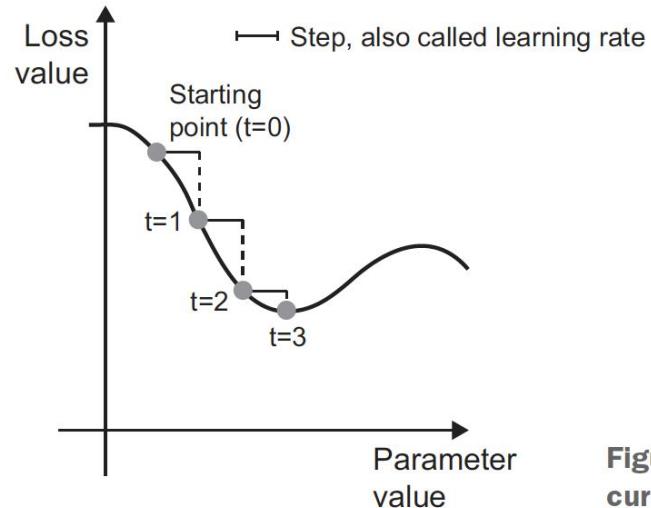


Figure 2.11 SGD down a 1D loss curve (one learnable parameter)

[Image source](#)



Stochastic Gradient Descent

- Easy to get stuck in local minima

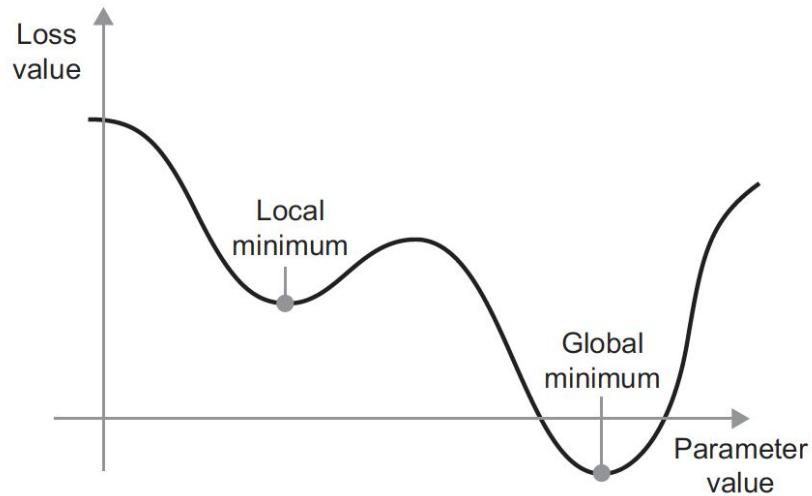


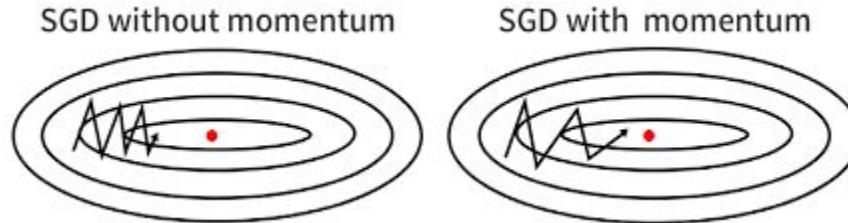
Figure 2.13 A local minimum and a global minimum

[Image source](#)



Momentum

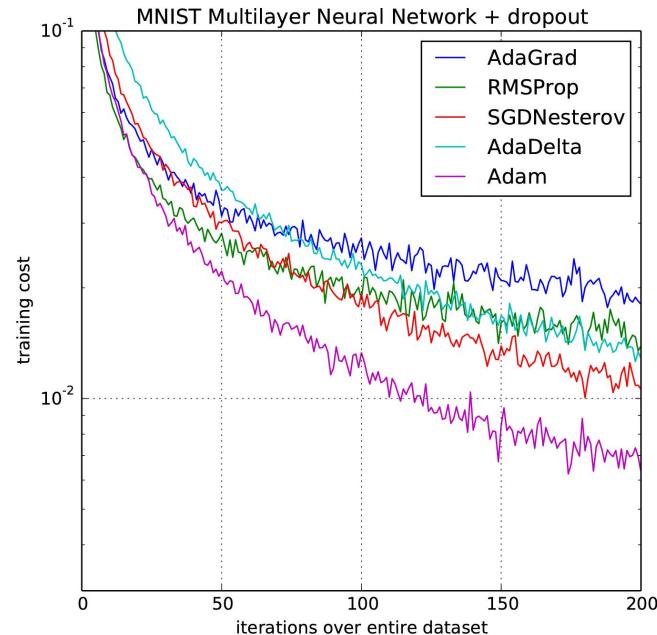
- Accelerates SGD: Like pushing a ball down a hill
- Take average of direction we've been heading (current velocity and acceleration)
- Limits oscillating back and forth, gets out of local minima



[Image source](#)

ADAM (ADAptive Moment Estimation)

Recommended default choice



[Image source](#)



ADAM: Momentum + RMSProp

- Momentum
 - Uses a moving average of the gradient
- RMSProp
 - Divides the learning rate element-wise by the root of the moving average of the squared gradient
 - More uniform learning steps, small gradients are impactful, faster convergence

$$w_t = w_{t-1} - \frac{\text{learning rate}}{\sqrt{\text{moving average of squared gradient}}} * \text{moving average of gradient}$$

[Image source](#)



DEMO: KERAS

Notebook: 01 Keras Fundamentals → 01.1-Linear Regression + 01.2-Keras

LAB: KERAS

Notebook: 01-Keras Fundamentals → 01.2L-Keras Lab

DEMO: CALLBACKS

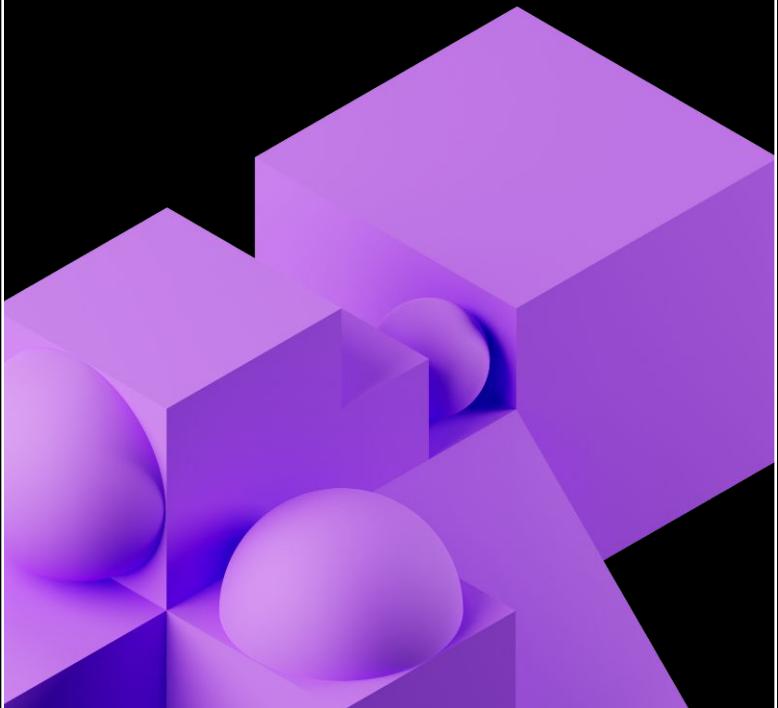
Notebook: 01 Keras Fundamentals → 01.4-Callbacks



LAB: CALLBACKS

Notebook: O1 Keras Fundamentals → O1.4L – Callbacks Lab

MLflow



Core Machine Learning Issues

Modern ML lifecycle comes with many challenges

- Keeping track of experiments or model development
- Reproducing code
- Comparing models
- Standardization of packaging and deploying models

MLflow addresses these issues.

MLflow Components

The four components of MLflow



Record and query experiments: code, data, config, results



Packaging format for reproducible runs on any platform



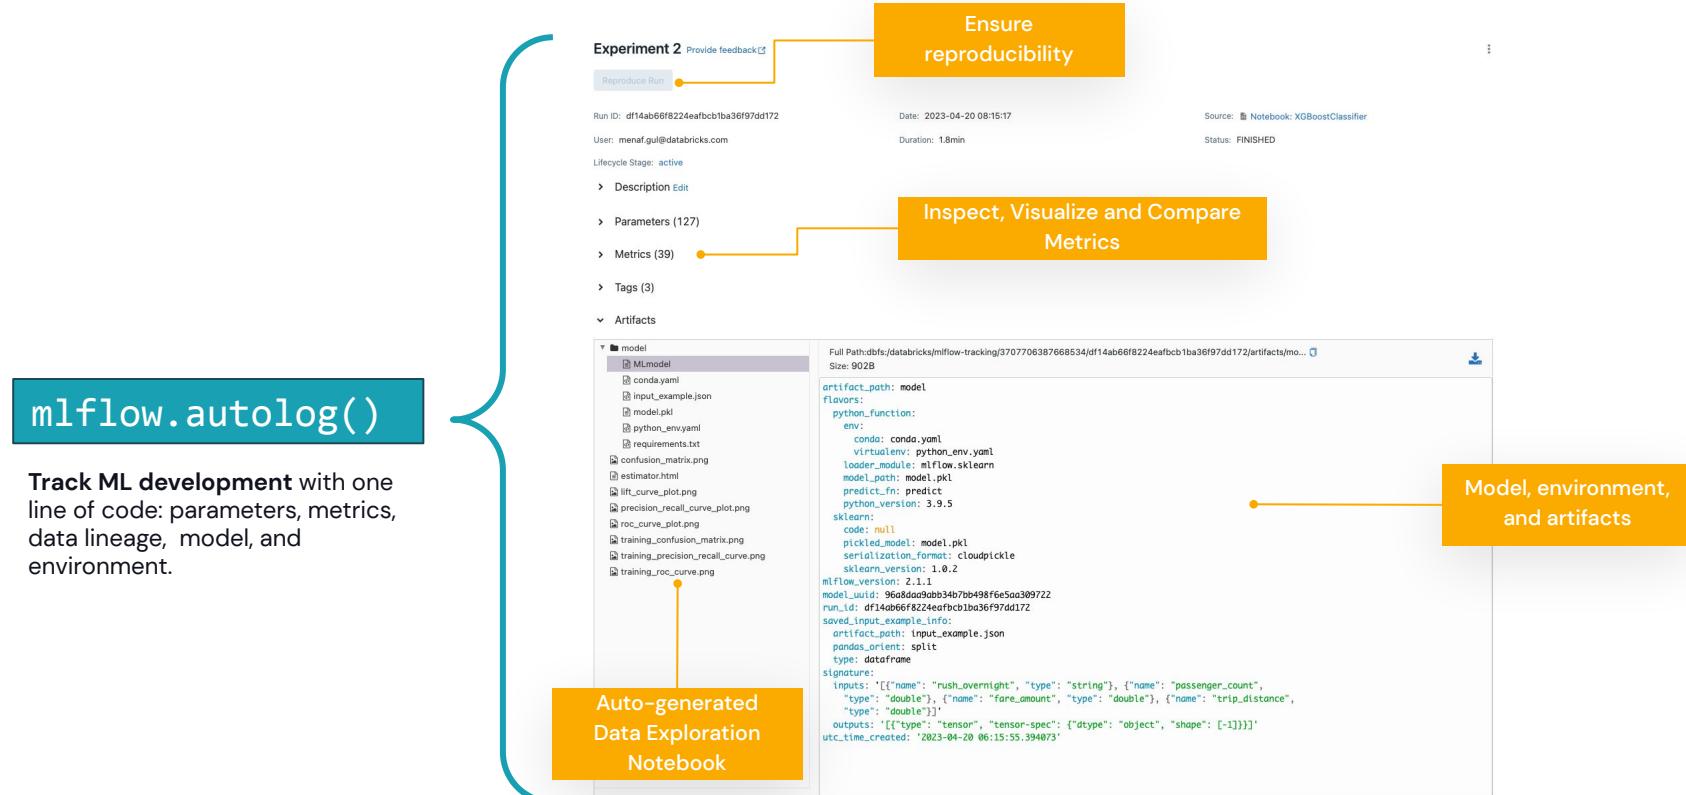
General model format that supports diverse deployment tools



Centralized and collaborative model lifecycle management

APIs: CLI, Python, R, Java, REST

MLflow Tracking and Auto-logging



DEMO: MLFLOW

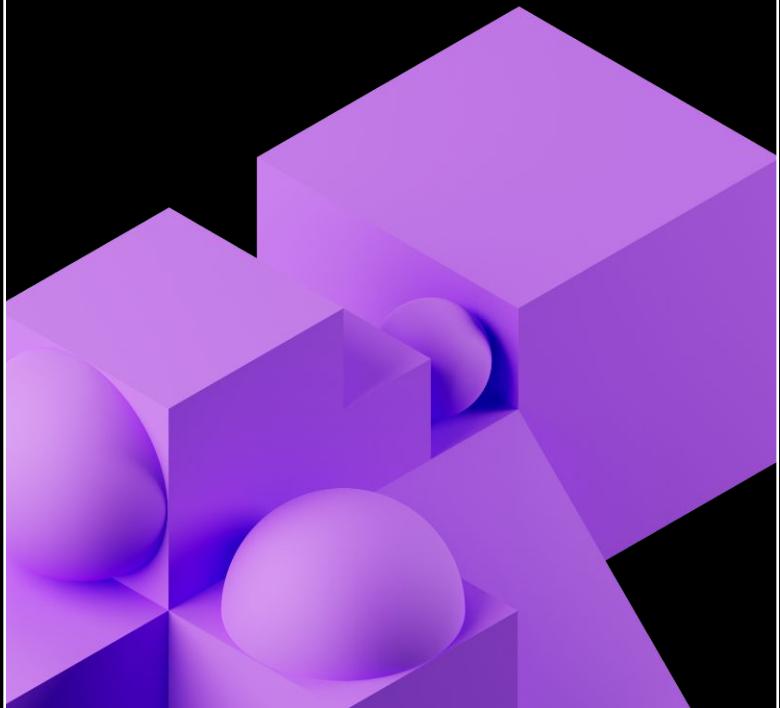
Notebook: O2-MLflow → O2.1-MLflow



LAB: MLFLOW

Notebook: O2-MLflow → O2.1L-MLflow Lab

Hyperparameter Tuning

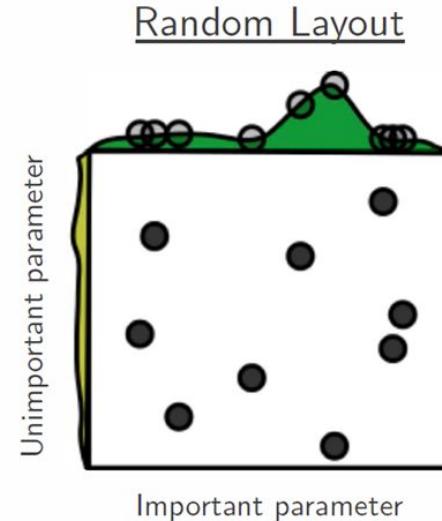
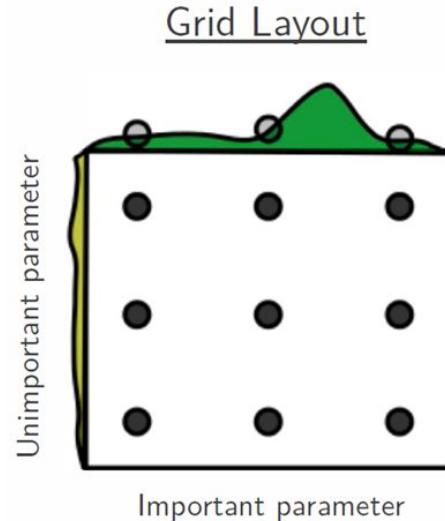


Problems with Grid Search

- Exhaustive enumeration is expensive
- Manually determined search space
- Past information on good hyperparameters isn't used
- So what do you do if...
 - You have a training budget
 - You have many hyperparameters to tune
 - You want to pick your hyperparameters based on past results

Optimizing Hyperparameter Values

Random Search



Random Search generally outperforms grid search

[Image source](#)



Hyperopt



- Open-source Python library
- Optimization over *awkward search spaces* (real-valued, discrete, and conditional dimensions)
- Supports serial or parallel optimization
- Spark integration
- Core algorithms for optimization:
 - Random Search
 - Adaptive Tree of Parzen Estimators (TPE)



DEMO: HYPEROPT

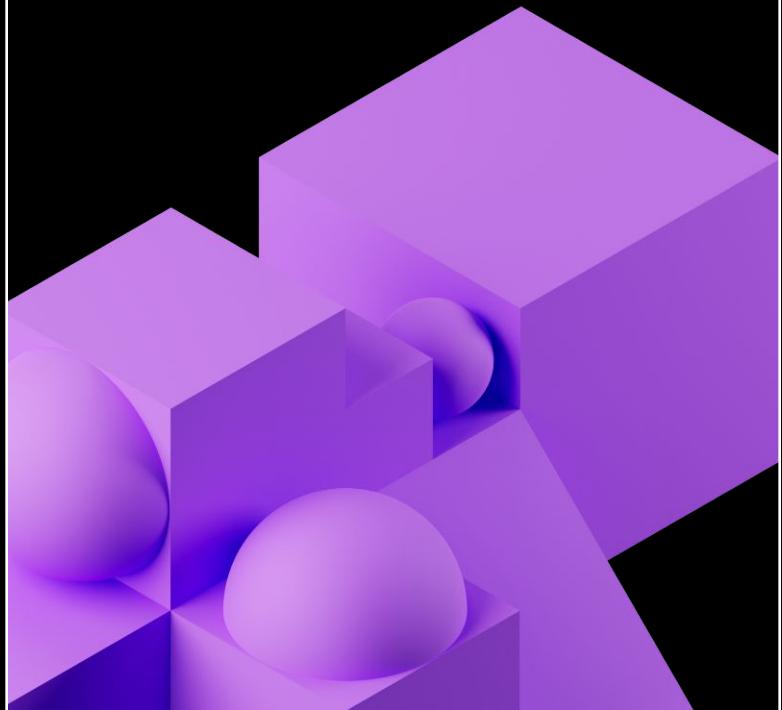
Notebook: 03-Hyperopt → 03.1-Hyperopt



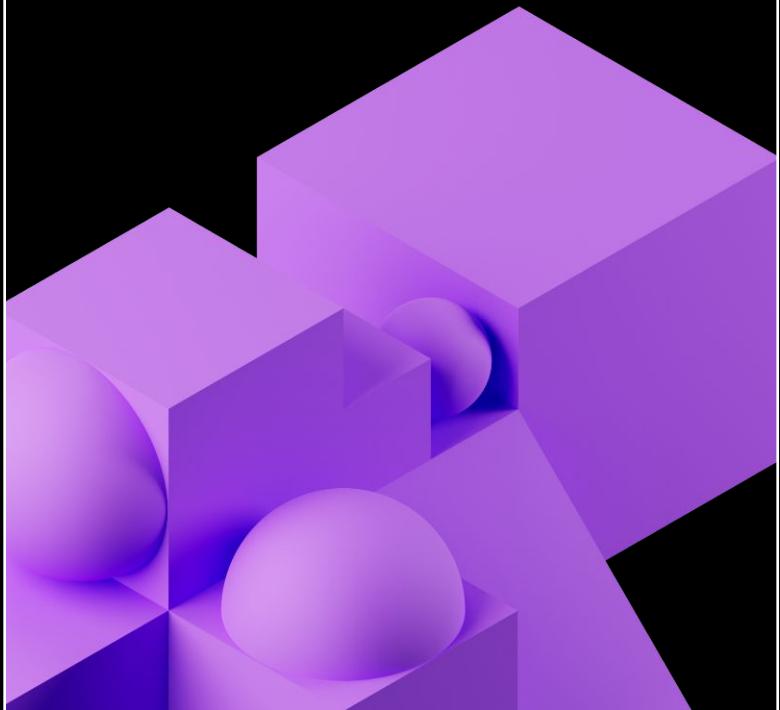
LAB: HYPEROPT

Notebook: 03-Hyperopt → 03.1L-Hyperopt Lab

Petastorm



Horovod



Horovod

Like this Russian dance, everybody dances in a circle and each person only touches their neighbor.



[Image source](#)

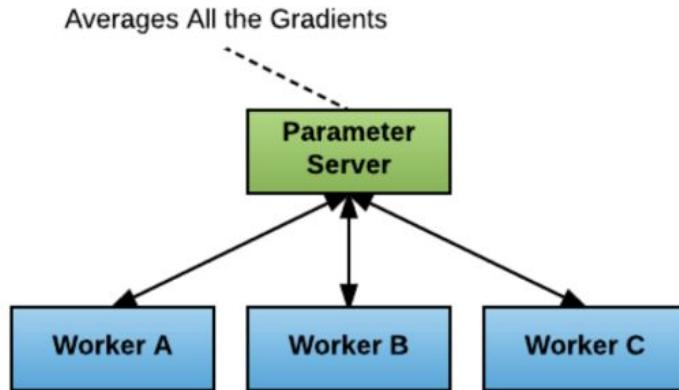


Horovod

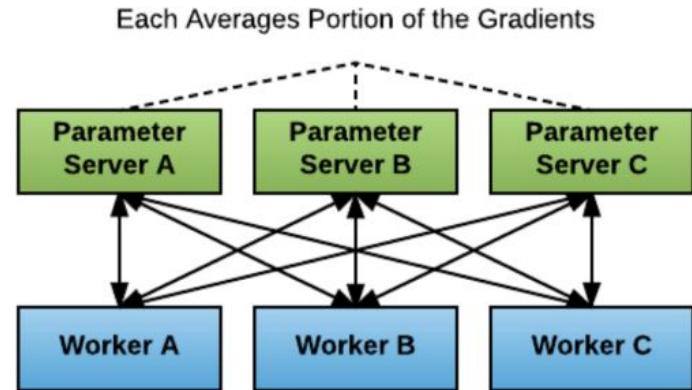
- Open-sourced by Uber in 2017
- Simplifies distributed neural network training
- Supports TensorFlow, Keras, PyTorch, and Apache MXNet



Classical Parameter Server



or

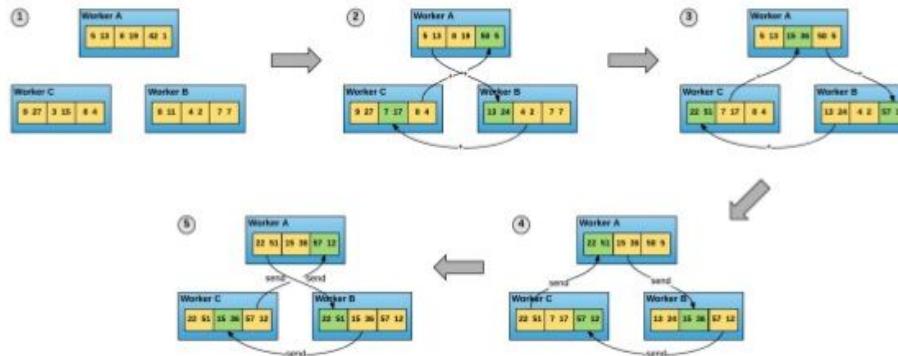


[Image source](#)



Ring All-Reduce

Horovod Technique



Patarasuk, P., & Yuan, X. (2009). Bandwidth optimal all-reduce algorithms for clusters of workstations. *Journal of Parallel and Distributed Computing*, 69(2), 117-124. doi:10.1016/j.jpdc.2008.09.002

UBER

Navigation icons

```
# Only one line of code change!
optimizer = hvd.DistributedOptimizer(optimizer)
```

[Image source](#)



DEMO: PETASTROM

Notebook: 04-Petastrom and Horovod → 04.1-Petastrom for Large Data



DEMO: HOROVOD

Notebook: 04-Petastrom and Horovod → 04.2-Horovod for Distributed Training

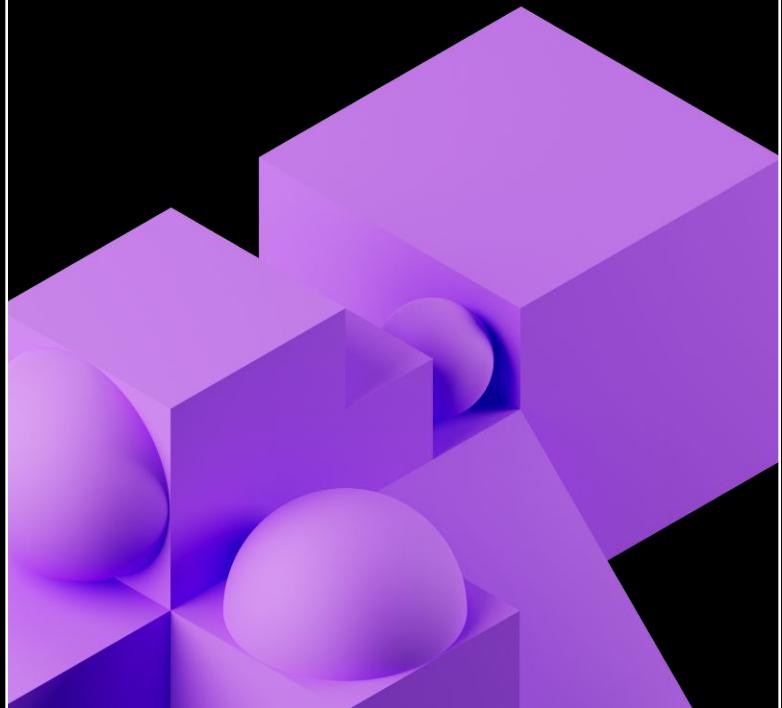


LAB: HOROVOD

Notebook: O4-Petastrom and Horovod → O4.2L-Horovod with Petastrom Lab



Convolutional Neural Networks

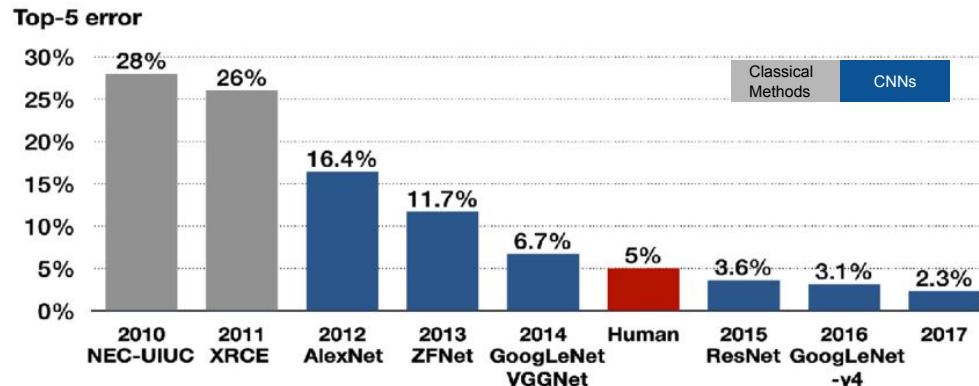




[Image source](#)

ImageNet Challenge

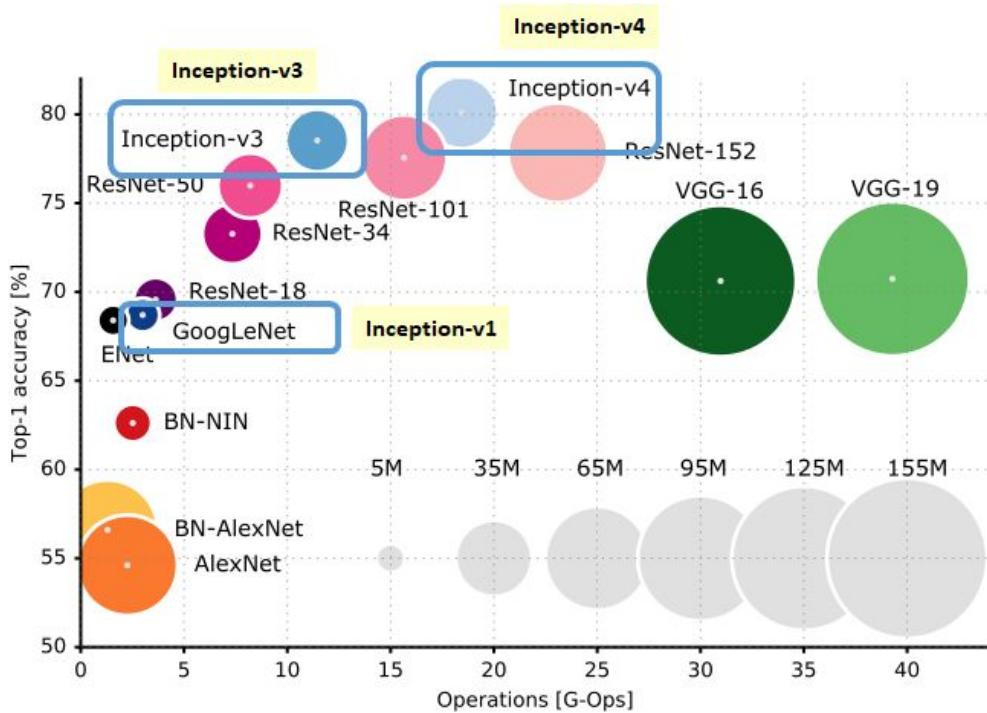
- Classify images in one of 1000 categories
- 2012 Deep Learning breakthrough with AlexNet: 16% top-5 test error rate (next closest was 25%)



[Image source](#)

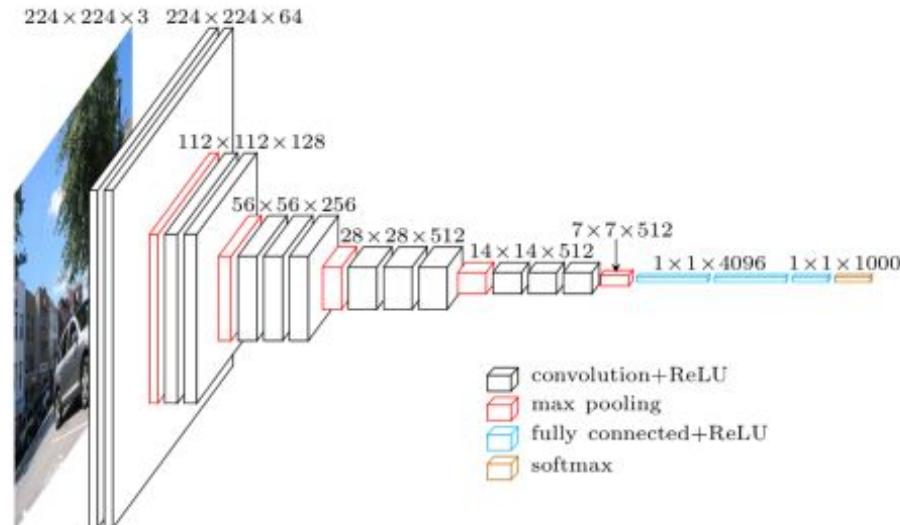


Top-1 Accuracy against Number of Operations



VGG16 (2014)

- One of the most widely used architectures for its simplicity



Convolutions

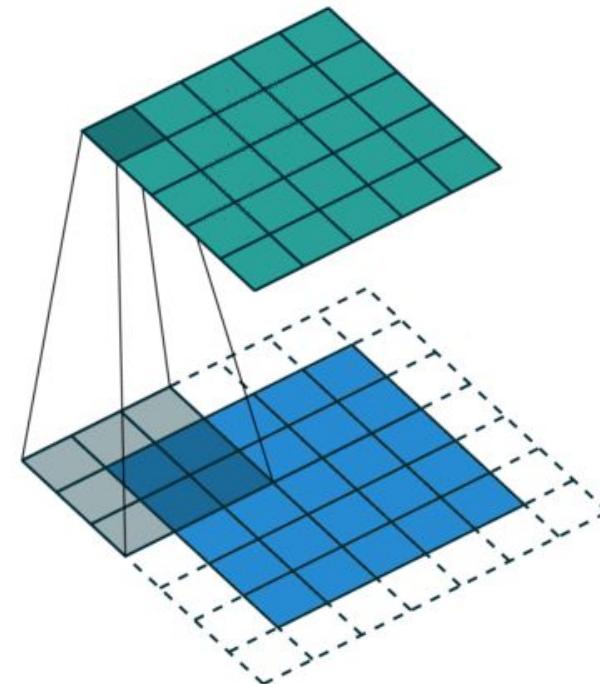
- Focus on spatial correlation (fewer parameters to learn)
- Filter/kernel slides across input image (often 3x3 and 5x5)

[Image Kernels Visualization](#)

[CS 231 Convolutional Networks](#)

Convolution

- Kernel size: Field of view
- Stride: Step size of kernel
- Padding: How to handle border
 - Valid: No padding (crops)
 - Same: Pads it so $\text{input_size} = \text{output_size}$



[Image source](#)



Max Vs Avg. Pooling

Max Pooling			
29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2
pool size

100	184
12	45

Average Pooling

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

2 x 2
pool size

36	80
12	15



VGG-16 Summary

Compute block1_conv1:

$$3 \times 3 \times 3 \times 64 + 64 = 1,792 \text{ params}$$

How to compute block1_conv2?

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 224, 224, 3)	0	
block1_conv1 (Convolution2D)	(None, 224, 224, 64)	1792	input_1[0][0]
block1_conv2 (Convolution2D)	(None, 224, 224, 64)	36928	block1_conv1[0][0]
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0	block1_conv2[0][0]
block2_conv1 (Convolution2D)	(None, 112, 112, 128)	73856	block1_pool[0][0]
block2_conv2 (Convolution2D)	(None, 112, 112, 128)	147584	block2_conv1[0][0]
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0	block2_conv2[0][0]
block3_conv1 (Convolution2D)	(None, 56, 56, 256)	295168	block2_pool[0][0]
block3_conv2 (Convolution2D)	(None, 56, 56, 256)	590080	block3_conv1[0][0]
block3_conv3 (Convolution2D)	(None, 56, 56, 256)	590080	block3_conv2[0][0]
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0	block3_conv3[0][0]
block4_conv1 (Convolution2D)	(None, 28, 28, 512)	1180160	block3_pool[0][0]
block4_conv2 (Convolution2D)	(None, 28, 28, 512)	2359808	block4_conv1[0][0]
block4_conv3 (Convolution2D)	(None, 28, 28, 512)	2359808	block4_conv2[0][0]
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0	block4_conv3[0][0]
block5_conv1 (Convolution2D)	(None, 14, 14, 512)	2359808	block4_pool[0][0]
block5_conv2 (Convolution2D)	(None, 14, 14, 512)	2359808	block5_conv1[0][0]
block5_conv3 (Convolution2D)	(None, 14, 14, 512)	2359808	block5_conv2[0][0]
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0	block5_conv3[0][0]
flatten (Flatten)	(None, 25088)	0	block5_pool[0][0]
fc1 (Dense)	(None, 4096)	102764544	flatten[0][0]
fc2 (Dense)	(None, 4096)	16781312	fc1[0][0]
predictions (Dense)	(None, 1000)	4097000	fc2[0][0]
Total params: 138357544			

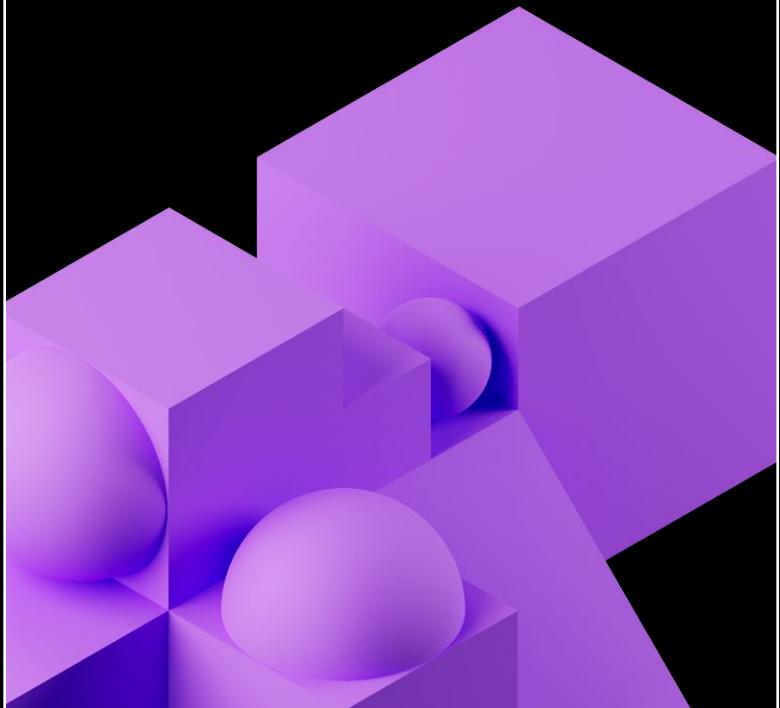


DEMO: CNN

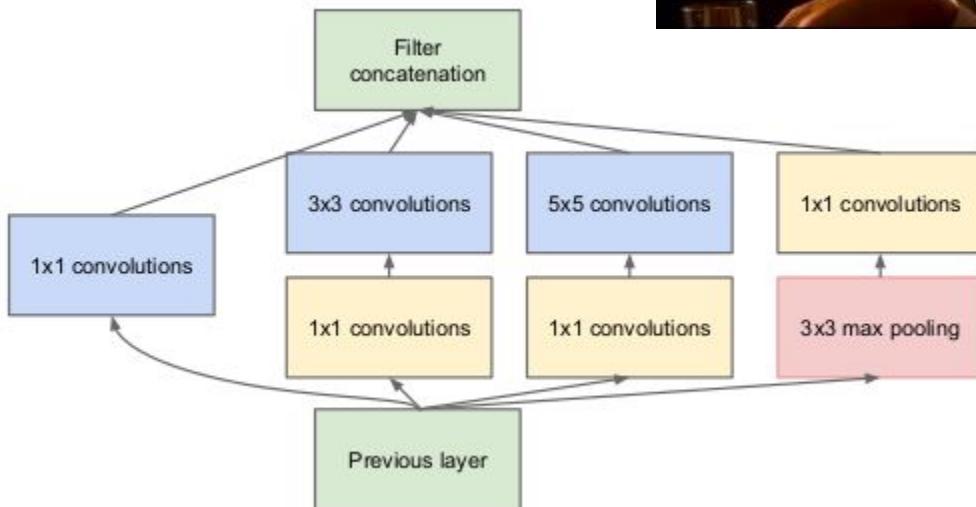
Notebook: 06-Convolutional Neural Networks → 06.1-Distributed Inference with CNNs



Evolution of CNN Architectures



GoogLeNet/Inception - 2014



GoogLeNet/Inception

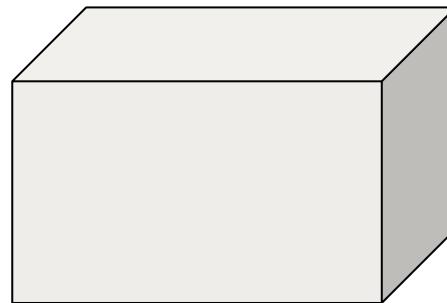
- Uses 1X1 convolutions to reduce dimension of input channels, then applies larger kernel
- Inception module:
 - Different sizes/types of convolutions for same input (effectively learns different features)
 - Stacks all the outputs
- Replaces the fully-connected layers at the end with global average pooling

Without 1x1 Convolution



14x14x512

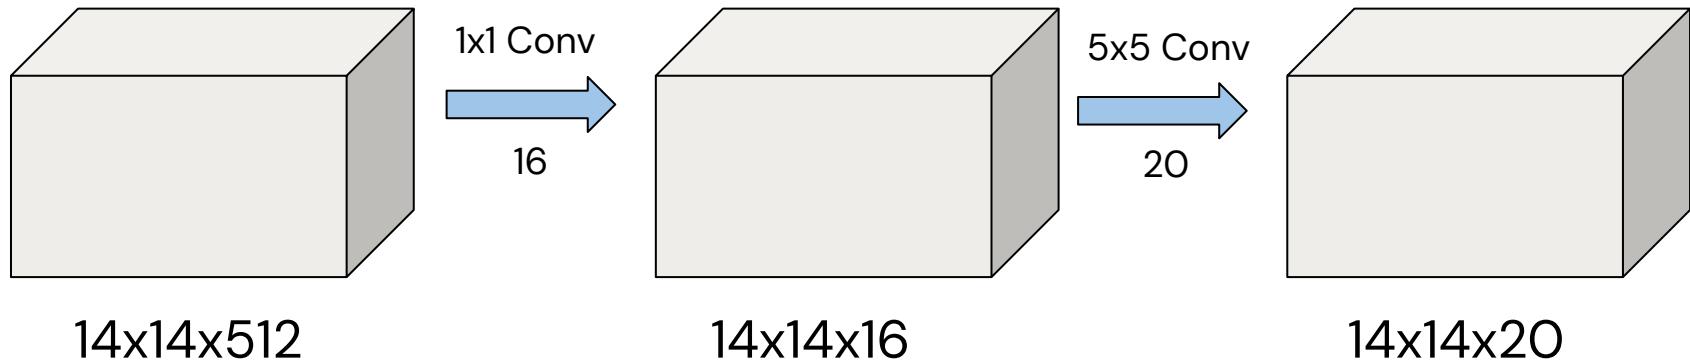
5x5 Conv
20



14x14x20

Number of operations: $(14 \times 14 \times 512) \times (5 \times 5 \times 20) = 50,176,000$

With 1x1 Convolution



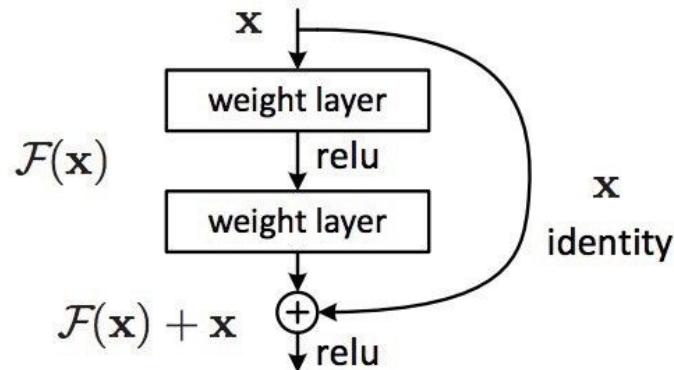
Number of operations for 1x1: $(14 \times 14 \times 512) \times (1 \times 1 \times 16) = 1,605,632$

Number of operations for 5x5: $(14 \times 14 \times 16) \times (5 \times 5 \times 20) = 1,568,000$

Total number of operations: 3,173,632 – just 6.3% of original parameters!

ResNet (Residual Network) – 2015

“Direct mappings are hard to learn...instead of trying to learn an underlying mapping from x to $H(x)$, learn the difference between the two, or the “residual.”



There are many other newer models

Network	Year	Top-5 ImageNet Accuracy	# of Params	Floating Point Operations
AlexNet	2012	84.7%	62M	1.5B
VGGNet	2014	92.3%	138M	19.6B
Inception v1	2014	93.3%	6.4M	2B
ResNet-152	2015	95.5%	60.3M	11B
Inception v3	2015	94.4%	23.8M	-
Xception	2016	94.5%	22.8M	-
NasNet	2017	95.3%	22.6M	-
MobileNet	2017	89.5%	4.24M	-
EfficientNet B5	2019	96.7%	30M	9.9B

Larger, Newer Models != Better

- MIT and Amazon researchers found ~3.4% errors across 10 benchmarking datasets. Label errors documented here:
<https://labelerrors.com>



MNIST given label:

8

We guessed: **9**

MTurk consensus: **9**

ID: 947



ImageNet given label:

tick

We guessed: **yellow garden spider**

MTurk consensus: **yellow garden spider**

ID: 00004095



CIFAR-10 given label:

airplane

We guessed: **ship**

MTurk consensus: **ship**



ID: 1718

Smaller models outperform on corrected labels

- Rankings of small models pre-trained on ImageNet increased
 - NasNet ranking drops from 1/34 → 29/34
 - Xception drops from 2/34 to 24/34
 - ResNet-18 increases from 34/34 → 1/34
 - ResNet-50 increases from 20/34 → 2/34
- Same trend occurs on the 13 models pre-trained on CIFAR-10
 - VGG-11 > VGG-19

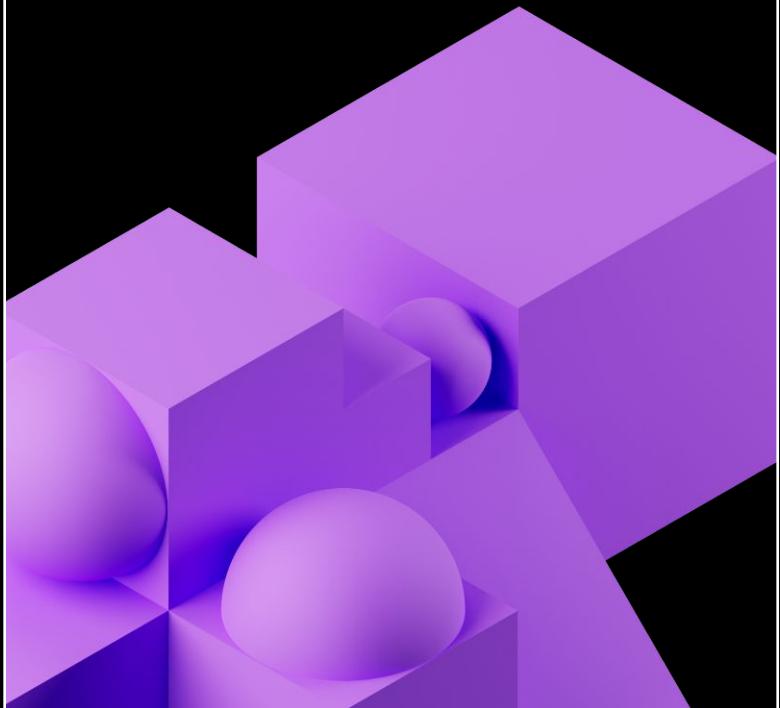
Network	Year	Top-5 ImageNet Accuracy	# of Params	Floating Point Operations
AlexNet	2012	84.7%	62M	1.5B
VGGNet	2014	92.3%	138M	19.6B
Inception v1	2014	93.3%	6.4M	2B
ResNet-152	2015	95.5%	60.3M	11B
Inception v3	2015	94.4%	23.8M	-
Xception	2016	94.5%	22.8M	-
NasNet	2017	95.3%	22.6M	-
MobileNet	2017	89.5%	4.24M	-
EfficientNet B5	2019	96.7%	30M	9.9B



Detrimental impact of errors increases with model size

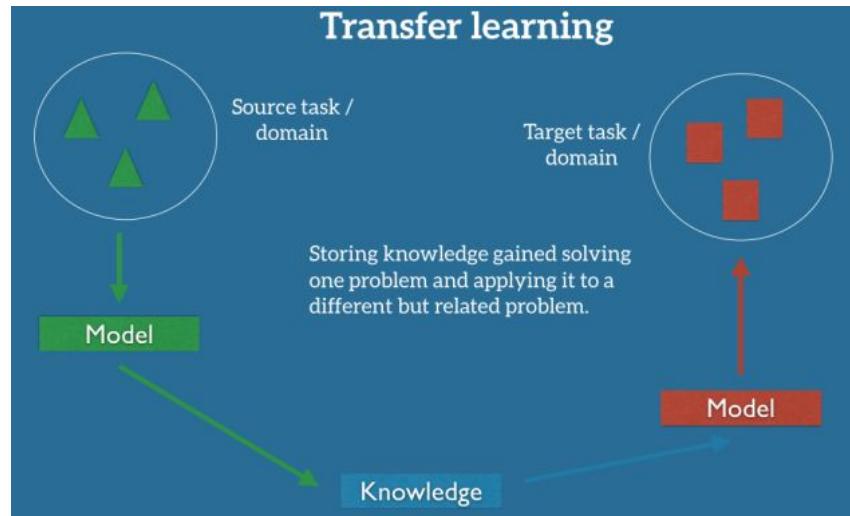
- Lower-capacity models
 - Provide regularization benefits
 - Are more resistant to learning asymmetric distribution of noisy labels
- Large models overfit to
 - specific benchmarks
 - quirks of the original label annotators
- Implications:
 - We only see the original test accuracy
 - But we should pick our models based on the corrected test accuracy
 - **SOTA on research data != SOTA on real production data**

Transfer Learning



Transfer Learning

- IDEA: Intermediate representations learned for one task may be useful for other related tasks



[Image source](#)



When To Use Transfer Learning?

	Similar dataset	Different dataset
Small dataset	Transfer learning: highest level features + classifier	Transfer learning: lower level features + classifier
Large dataset	Fine-tune*	Fine-tune*

[Reference: Andrej Karpathy's Transfer Learning](#)



DEMO: TRANSFER LEARNING

Notebook: 07-Transfer Learning → 07.1-Transfer Learning with Data Generators

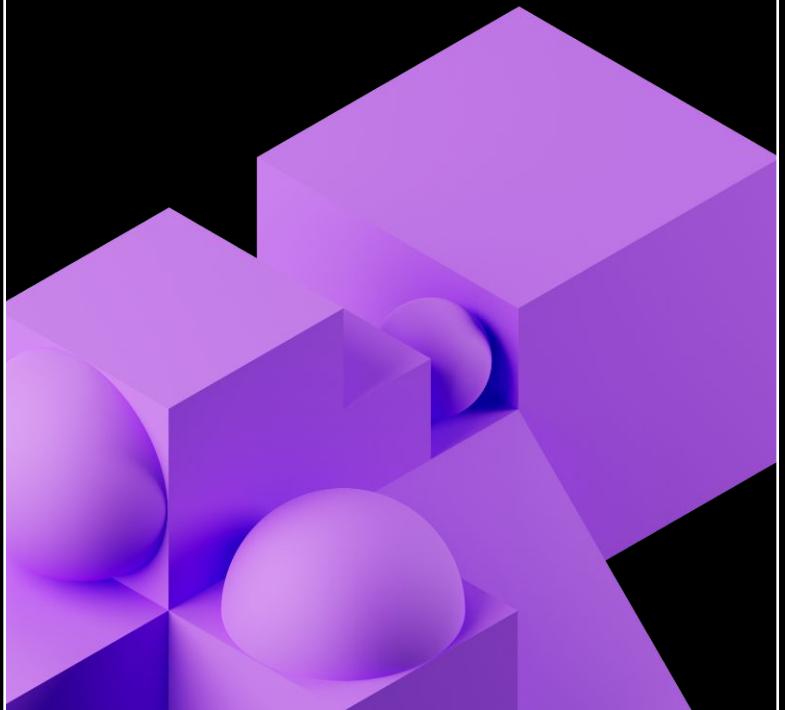


LAB: TRANSFER LEARNING

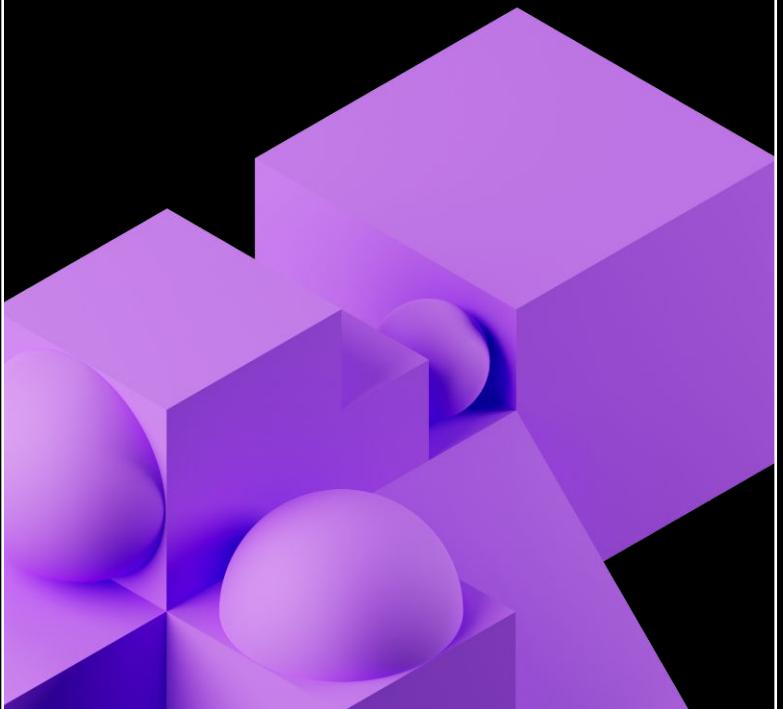
Notebook: 07-Transfer Learning → 07.2L – Transfer Learning for CNNs Lab



Questions?



Summary and Next Steps

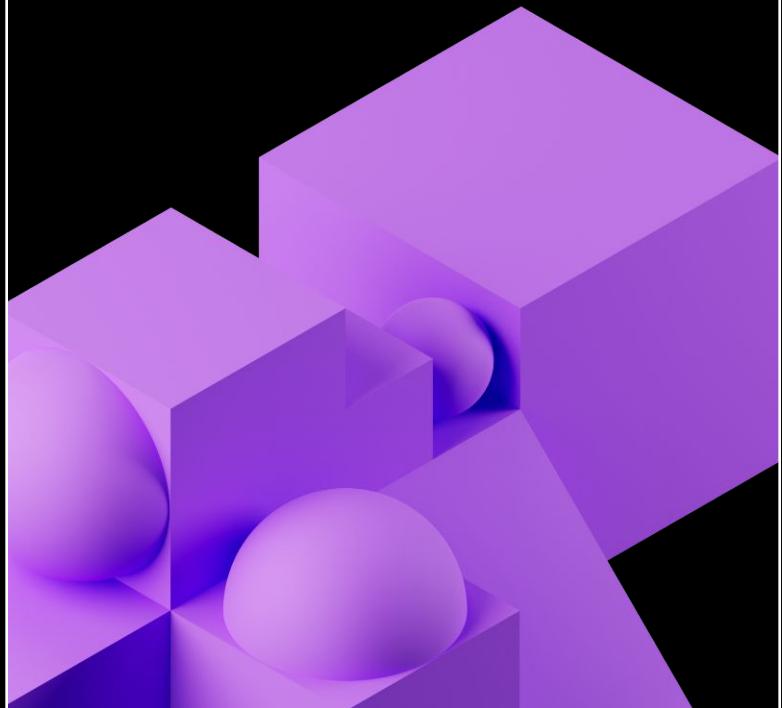


THANK YOU!



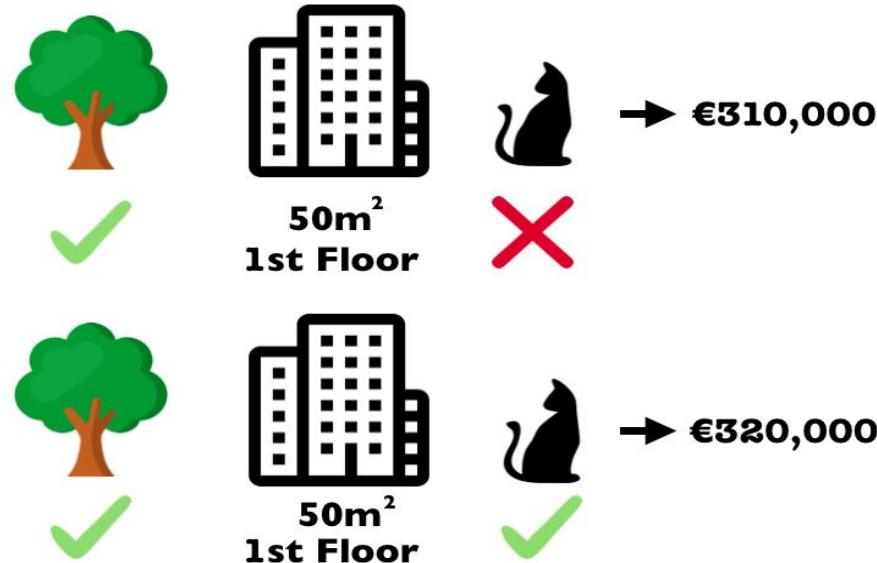
Appendix

Model Interpretability



SHAP (SHapley Additive exPlanations)

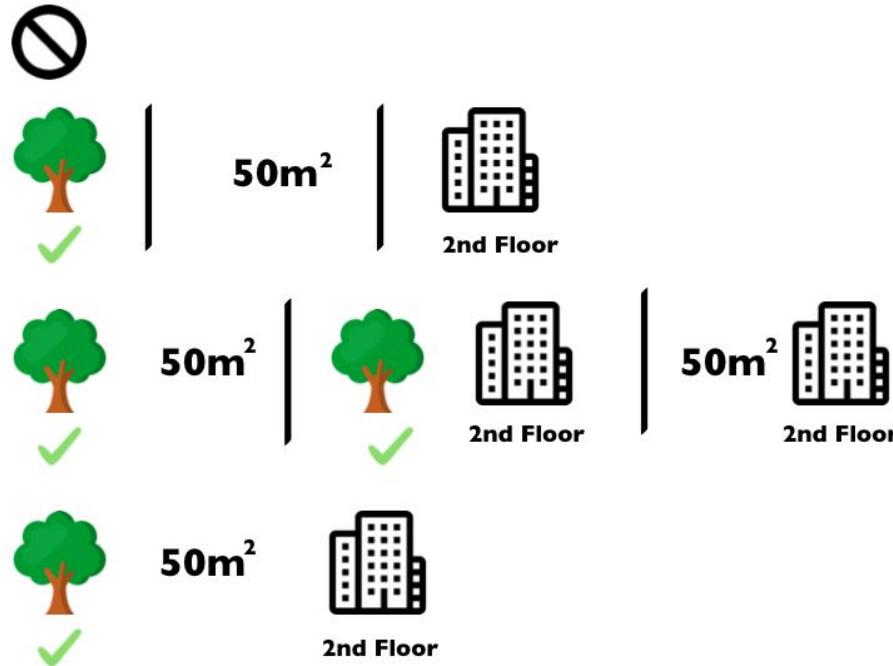
- Shapley Values



[Image source](#)



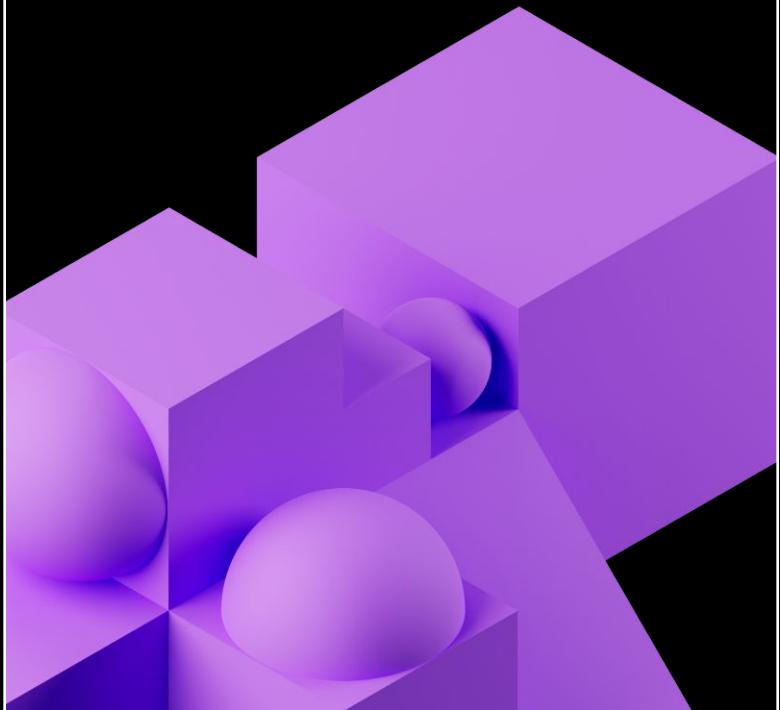
SHAP



[Image source](#)



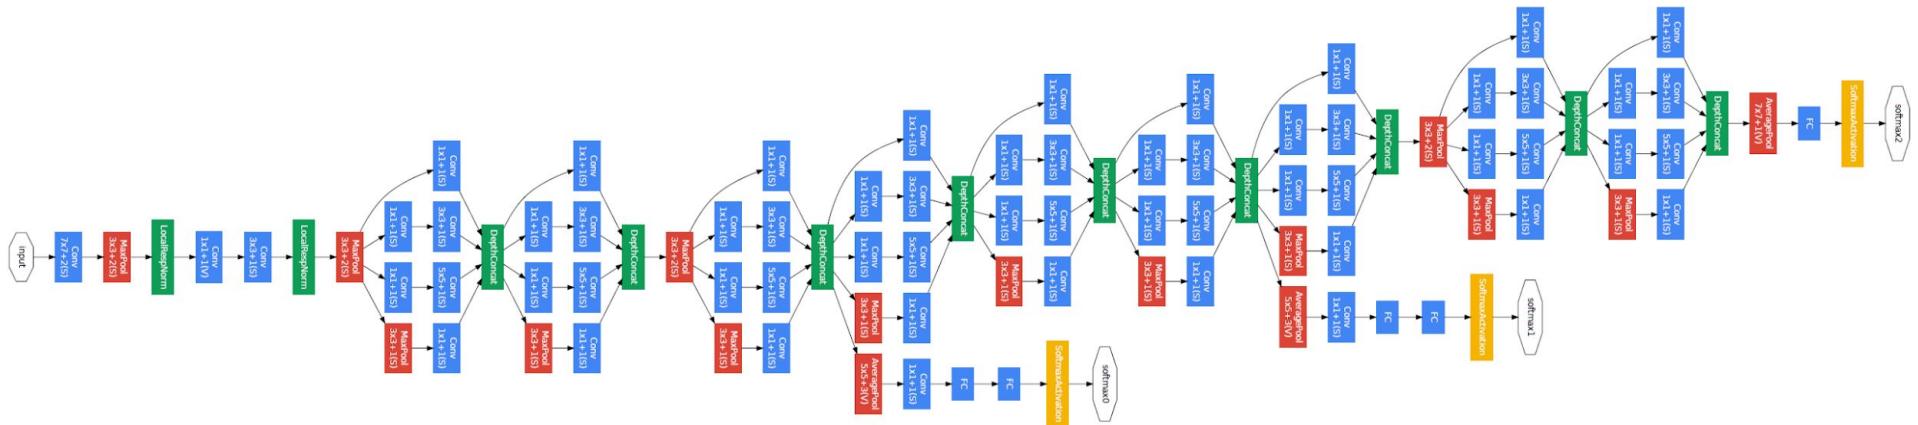
CV



What Do CNNs Learn?

- Breaking Convnets

Inception Model: 22 Layers in Total



Testing Details

- 7 models ensemble
 - 1 has a different architecture, how are the other 6 different?
- Multi-scale testing: 256, 288, 320, and 352 dimensions (4 scales)
- Multi-crop testing
- 4 scales×3 squares×6 crops×2 mirrored versions=144 crops/image

Number of models	Number of Crops	Cost	Top-5 error	compared to base
1	1	1	10.07%	base
1	10	10	9.15%	-0.92%
1	144	144	7.89%	-2.18%
7	1	7	8.09%	-1.98%
7	10	70	7.62%	-2.45%
7	144	1008	6.67%	-3.45%

Inception v2/BN-Inception

- Adds Batch Norm
 - Normalizes layer inputs, allows for higher learning rate
- Replaces 5×5 conv with two 3×3 convs to reduce params

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;
Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

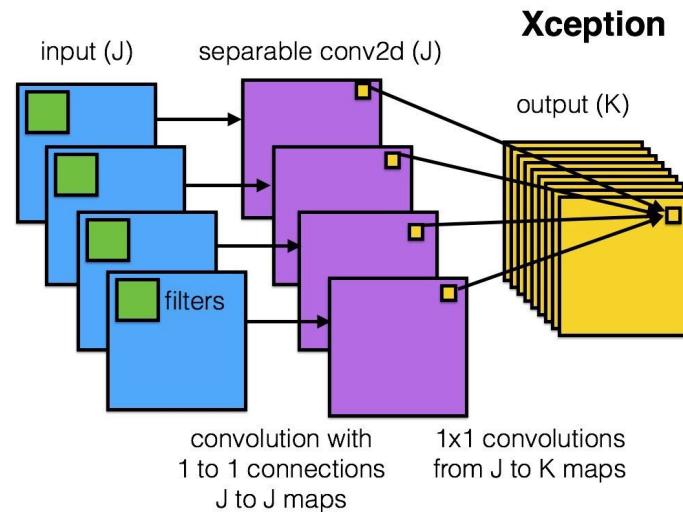
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$



Xception (Extreme Inception) – 2017

“Cross-channel correlations and spatial correlations are sufficiently decoupled that it is preferable not to map them jointly.”



Where is CV headed?

- We've overfit to
ImageNet/MNIST/CIFAR...
 - [COCO](#) (MSFT/Facebook)
 - [Open Images Dataset](#) (Google)

What is COCO?



COCO is a large-scale object detection, segmentation, and captioning dataset. COCO has several features:

- ✓ Object segmentation
- ✓ Recognition in context
- ✓ Superpixel stuff segmentation
- ✓ 330K images (>200K labeled)
- ✓ 1.5 million object instances
- ✓ 80 object categories
- ✓ 91 stuff categories
- ✓ 5 captions per image
- ✓ 250,000 people with keypoints



databricks